

Jianming Tong

+1 470-357-8082 | jianming.Tong@gatech.edu | [jianmingTong.github.io](https://github.com/jianmingTong)

RESEARCH INTEREST

Computer Architecture; Hardware-software co-design; Deep Learning Compiler and Accelerator; On-chip Network; FPGA; VLSI Design

EDUCATION

B.S. in Electrical Engineering, Sep 2016 – June 2020

Xi'an JiaoTong University (XJTU), Xi'an

- Advisor: Dr. [Pengju Ren](#)
- GPA: 3.91/4.3, Rank: 6/361

Visiting Ph.D. Student, July 2020 – Jan 2021

Tsinghua University (THU), Beijing

- Advisor: Dr. [Yu Wang](#)

Ph.D. in Computer Science, Jan 2021 – Dec 2025

Georgia Institute of Technology (GT), Atlanta

- Advisor: Dr. [Tushar Krishna](#)
- GPA: 4.0/4.0

PEER-REVIEWED PUBLICATIONS (* EQUAL CONTRIBUTION)

- **A Reconfigurable Accelerator with Data Reordering Support for Low-Cost On-Chip Dataflow Switching**
Jianming Tong, Anirudh Itagi, Tushar Krishna.
3rd On-Device Intelligence Workshop, @ In Proc of Sixth Conference on Machine Learning and Systems (MLSys), Jun 2023
- **ReLU-FHE: Low-cost Accurate ReLU Polynomial Approximation in Fully Homomorphic Encryption Based ML Inference**
Jingtian Dang*, **Jianming Tong***, Anupam Golder, Callie Hao, Tushar Krishna.
3rd On-Device Intelligence Workshop, @ In Proc of Sixth Conference on Machine Learning and Systems (MLSys), Jun 2023
- **SUSHI: SUBgraph Stationary Hardware-software Inference Co-design**
Payman Behnam*, **Jianming Tong***, Alind Khare, Yangyu Chen, Yue Pan, Pranav Gadikar, Abhimanyu Rajeshkumar Bambhaniya, Tushar Krishna, Alexey Tumanov.
In Proc of Sixth Conference on Machine Learning and Systems (MLSys), Jun 2023
- **FPGA-Based High-Performance Real-Time Emulation of Radar System using Direct Path Compute Model**
Xiangyu Mao*, Mandovi Mukherjee*, Nael M. Rahman*, Uday Kamal, Sudarshan Sharma, Payman Behnam, **Jianming Tong**, Jongseok Woo, Coleman B DeLude, Joseph W. Driscoll, Jamin Seo, Santosh Pande, Tushar Krishna, Justin Romberg, Madhavan Swaminathan, and Saibal Mukhopadhyay
In Proc of IEEE MTT-S International Microwave Symposium (IMS), Jun 2023
- **A High Performance Computing Architecture for Real-Time Digital Emulation of RF Interactions**
Mandovi Mukherjee*, Nael Mizanur Rahman*, Coleman B. DeLude*, Joseph W. Driscoll*, Uday Kamal, Jongseok Woo, Jamin Seo, Sudarshan Sharma, Xiangyu Mao, Payman Behnam, Sharjeel M. Khan, Daehyun Kim, Jianming Tong, Prachi Sinha, Santosh Pande, Tushar Krishna, Justin Romberg, Madhavan Swaminathan, and Saibal Mukhopadhyay
In Proc of IEEE Radar Conference (RadarConf), May 2023
- **FastSwitch: Enabling Real-time DNN Switching via Weight-Sharing**
Jianming Tong, Yangyu Chen, Yue Pan, Abhimanyu Bambhaniya, Alind Khare, Taekyung Heo, Alexey Tumanov, and Tushar Krishna.
The 2nd Architecture, Compiler, and System Support for Multi-model DNN Workloads Workshop @ISCA, 2022.
- **A Configurable Architecture for Efficient Sparse FIR Computation in Real-time Radio Frequency Systems**
Jamin Seo, Nael Mizanur Rahman, Mandovi Mukherjee, Coleman DeLude, **Jianming Tong**, Justin Romberg, Tushar Krishna, and Saibal Mukhopadhyay.
IEEE Microwave and Wireless Technology Letters (IMS), 2022.

- **SMMR-Explore: SubMap-based Multi-Robot Exploration System with Multi-robot Multi-target Potential Field Exploration Method**
Jincheng Yu*, **Jianming Tong***, Yuanfan Xu, Zhilin Xu, Haolin Dong, Tianxiang Yang and Yu Wang
2021 IEEE International Conference on Robotics and Automation (ICRA 2021, oral) [[code](#)], [[Demo Link](#)].
- **ac2SLAM: FPGA Accelerated High-Accuracy SLAM with Heapsort and Parallel Keypoint Extractor**
Cheng Wang, Yinkun Liu, Kedai Zuo, **Jianming Tong [Project Leader]**, Yan Ding, and Pengju Ren.
International Conference on Field-Programmable Technology (FPT 2021, Full Paper) [[code](#)].
- **PIT: Processing-In-Transmission with Fine-Grained Data Manipulation Networks**
Tian Xia, Pengchen Zong, Haoran Zhao, **Jianming Tong**, Wenzhe Zhao, Nanning Zheng and Pengju Ren
IEEE Transactions on Computers (TOC)
- **COCOA: Content-Oriented Configurable Architecture based on Highly-Adaptive Data Transmission Networks**
Tian Xia, Pengchen Zong, Haoran Zhao, **Jianming Tong**, Wenzhe Zhao, Nanning Zheng and Pengju Ren
The 30th edition of the ACM Great Lakes Symposium on VLSI (GLSVLSI 2020)

Books

- **On Chip Networks, Second Edition [Translated Book]**
Natalie Enright Jerger, Tushar Krishna, and Li-Shiuan Peh
Synthesis Lectures on Computer Architecture, Morgan Claypool Publishers, June 2017
Chinese Translator: Pengju Ren, Tian Xia, **Jianming Tong [Project Leader]**, Pengcheng Zong, Haoran Zhao.
Publishing House of Electronics Industry, Jan 2021

AWARD

Finalist in Qualcomm Fellowship	Oct 2022, Apr 2023
2 nd Georgia Tech SCS Poster Competition	Apr 2022
National Encouragement Scholarship (top 5%)	Sep 2017, Sep 2018, Sep 2019
Huawei Scholarship (top 6/60)	May 2019

INDUSTRY EXPERIENCE

Performance Modeling Rivos Inc, CA

- Mentor: [Gautham Chinya](#); **Research Intern**, May 2023 – Aug 2023
- Dedicated on-chip network performance modeling (Stealth-mode company, no details disclosed)

End-to-end Framework for Inference Pacific Northwest National Laboratory, WA

- Mentor: [Roberto Gioiosa](#); **Research Intern**, Jun 2022 – Aug 2022
- Designed end-to-end inference framework on heterogeneous clusters with Xilinx VCK 5000, AMD CPU & GPU.
- Designed full-stack end-to-end inference including quantization, compiler, run-time and FPGA accelerator.

Fully Homomorphic Encryption (FHE) Acceleration Architecture Design DAMO Academy Alibaba Inc., Beijing

- Mentor: [Jiansong Zhang](#); **Research Intern**, Jun 2021 – Aug 2021
- Designed efficient homomorphic encryption acceleration architecture and deployed on Xilinx Alevo U280 for PoC.
- Categorized serial and parallel accelerator architectures for Number Theoretic Transform (NTT).

SKILLS

Programming	C/C++, Python, OpenCL, MCL, LLVM, Clang, MLIR, (System) Verilog, Xilinx HLS
Tools	GEM5, SST, Scale-Sim, MAESTRO, Timeloop, Xilinx Vivado, Vitis (AI), Cadence, Synapse

MENTORSHIP

Name	First employment
Yangyu Chen (Master@GT 2023)	Apple ASIC Verification Designer
Yue Pan (undergrade@GT 2022)	UCSD Ph.D.
Yuqi He (undergrade@GT 2022)	Apple ASIC Designer
Jingtian Dang (undergrade@GT 2022)	CMU ECE Master
Yingkun Liu (undergrade@XJTU 2021)	SJTU Ph.D.

Kedai Zuo (undergrade@XJTU 2021)
Cheng Wang (undergrade@XJTU 2021)

UCSD Master
Tsinghua Ph.D.

PROJECTS

MAERI 2.0: An end-to-end Inference Framework for enabling accelerator compiler research

Advisor: [Tushar Krishna](#) and [Alexey Tumanov](#), Georgia Tech

Jan 2021 – now

- Extended MAERI with memory-hierarchy, compiler and running time to verified end-to-end inference framework.
- Enabled “searched mapping from MAESTRO” deployment and verified on Xilinx VCK, Alevo and ZYNQ FPGA.
- Designed fully scalable, parameterized and verified architecture template using Xilinx HLS.
- <https://maeri-project.github.io/> (**Tutorial @ICS 2022**).

(IARPA) Large-scale Network-on-Chip (NoC) Modeling and Simulation in SST Simulator.

Advisor: [Tushar Krishna](#), Georgia Tech, collaborated with Intel.

Nov 2022 – now

- Modeling, simulating and testing large-scale multi-dimension NoC in [SST simulator toolkit](#).

(DARPA) High-performance Network-on-Chip (NoC) Design enabling arbitrary multicasting and unicasting

Advisor: [Tushar Krishna](#) and [Saibal Mukhopadhyay](#), Georgia Tech

Jan 2021 – Jan 2022

- Designed scalable NoC supporting arbitrary multicasting and unicasting from 1024 sources to 204 destinations.
- Synthesized, PNR and tapped out the designed NoC in scalable chip under TSMC 24 nm.

TorchFHE: An PyTorch library for privacy-preserving neural network inference using FHE.

Advisor: [Tushar Krishna](#) and [Callie Hao](#), Georgia Tech

Mar 2022 – now

- Enabled secure neural network inference in fully homomorphic encryption domain using BFV/CKKS schemes.

SELECTED TALKS

Enable Best ML Inference and Training: A systematic Approach @ EIC Lab – GaTech

Mar 2023

Full-Stack ML Dataflow, Mapping and SW/HW Co-Design and Search @ NICSEFC Lab – Tsinghua

Nov 2022