

Jianming Tong

+1 470-357-8082 | jianming.tong@gatech.edu | jianmingtong.github.io

RESEARCH INTEREST

System for AI and Cryptography

Enabling today's AI systems to work in a privacy-preserving manner without sacrificing performance

EDUCATION

- Jan'21 – now **Georgia Institute of Technology, PhD** in Computer Science
Advisor: [Tushar Krishna](#), Thesis: Enabling AI ASICs for Cryptography Primitives
- Jan'21 – May'24 **Georgia Institute of Technology, MS** in Computer Science
Advisor: [Tushar Krishna](#), Thesis: Reconfigurable Dataflow Accelerator [[ISCA'24](#)]
- Sep'16 – May'20 **Xi'an JiaoTong University, BE** in Electrical Engineering
Advisor: [Pengju Ren](#), Thesis: FPGA Accelerated High-Accuracy SLAM [[FPT'21](#)]

PROFESSIONAL EXPERIENCE

- Aug'24 – Apr'25 **Student Researcher in Google**, Cambridge, MA
Leveraging Google TPU for Homomorphic Encryption, host: [Asra Ali, Jevin Jiang](#)
- Sep'23 – Feb'25 **Visiting Researcher in Massachusetts Institute of Technology**, Cambridge, MA
Focus on full-stack acceleration for FHE ([HPCA'26](#)), host: [Arvind](#)
- Jan'24 – May'21 **Teaching Assistant in Massachusetts Institute of Technology**, Cambridge, MA
Constructive Computer Architecture (6.192) – Instructor: [Arvind, Tushar Krishna](#)
- Jan'21 – present **Graduate Research Assistant in Georgia Institute of Technology**, Atlanta, GA
Software-System-Hardware Codesign for Edge ML ([MLSys'23, IEEE MICRO'23](#))
- May'23 – Aug' 23 **Engineer Intern in Rivos Inc.** Mountain View, CA, Mentor: [Gautham Chinya](#)
Inter-chiplet performance modelling (NDA)
- Jan'23 – May'23 **Teaching Assistant in Georgia Institute of Technology**, Atlanta, GA
Processor Design (CS 3220) – Homework / Project Design
- Jun' 22 – Aug' 22 **Research Intern in Pacific Northwest National Laboratory**, Remote
A single-author end-to-end FPGA framework for AI inference ([Tutorial@ICS'22](#))
- Jul'21 – Aug'21 **Research Intern in DAMO Academy Alibaba Inc.** China, Mentor: [Jiansong Zhang](#)
Designed and implemented FPGA accelerator for Homomorphic Encryption scalar/vector multiplication and addition (open-sourced with paper [DAC'23](#))
- Aug'20 – Jan'21 **Visiting Ph.D. in Tsinghua University**, Beijing, China, Mentor: [Yu Wang](#)
Designed multi-robot collaborative exploration alg. (co-first author paper [ICRA'21](#))
- Sep'18 – Jul'20 **Undergraduate Research Assistant in Xi'an JiaoTong University**, Xi'an, China
Lead NoC book translation and robotic HW acceleration (open-sourced paper [FPT'21](#))

SELECTED AWARDS AND HONORS [\[Full list\]](#)

- Aug'25 **GT NEXT Award** (@[Geogia Tech](#))
- Jun'25 **2nd place University DEMO** @Design Automation Conference ([DAC'25](#))
- Jun'24 **ML and System Rising Star** @ [MLCommon](#)
- Apr'24 **DAC Young Fellow** @ Design Automation Conference ([DAC'24](#))
- Sep'23 **Best Poster Award** – **SUSHI** @ Industry-Academia Partner Workshop ([IAP'23](#))
The top voted poster among over 20+ candidates by industry partners.
- Jul'23 **Qualcomm Innovation Fellowship** – Latency/Accuracy Navigation in Edge ML
18 winners out of 182 submissions (three rounds, nationwide)
- Oct'22 **Qualcomm Innovation Fellowship Finalist** – ML Accel. Side-channel Attack

PUBLICATIONS (* EQUAL CONTRIBUTION) -- CONFERENCE

In Submission +Deployed in Google Cloud (TPU series)	Enabling Multi-scalar Multiplication over AI ASICs for ZKP Jianming Tong , Jingtian Dang, Simon Langowski, Tianhao Huang, Asra Ali, Jeremy Kun, Jevin Jiang, Srinivas Devadas, Tushar Krishna. [Code]
HPCA 2026 +Deployed in Google Cloud (TPU series) +2 nd place University DEMO@DAC'25 +GT NEXT Award	Leveraging ASIC AI Chips for Homomorphic Encryption Jianming Tong , Tianhao Huang, Leo De Castro, Anirudh Itagi, Jingtian Dang, Anupam Golder, Asra Ali, Jevin Jiang, Arvind, G. Edward Suh, Tushar Krishna. IEEE International Symposium on High-Performance Computer Architecture, 2026 [Code]
USENIX Security	Privatar: Enabling Privacy-preserving Real-time Multi-user VR through Secure Outsourcing Jianming Tong , Hanshen Xiao, Hao Kang, Krishnakumar Nair, Ashish Sirasao, G. Edward Suh, Tushar Krishna. USENIX Security Poster, Aug 2025 [Code]
HASP 2025	S ² Loop: Explore Optimal Authentication Block Strategy for ML Jan Strzeszynski*, Jianming Tong *, Kyungmi Lee, Nathan Xiong, Angshuman Parashar, Joel Emer, Tushar Krishna, Mengjia Yan Hardware and Architectural Support for Security and Privacy, Oct 2025 [Code]
ISPASS 2025	SCALE-Sim v3: A Modular Cycle-Accurate Systolic Accelerator Simulator for End-to-End System Analysis Ritik Raj, Sarbartha Banerjee*, Nikhil Srinivas*, Zishen Wan*, Jianming Tong *, Ananda Samajdar, Tushar Krishna. [Code]
ISPASS 2025	Constrained Dataflow Accelerator for Real-Time Multi-Task Multi-Model Machine Learning Workloads Jamin Seo, Jianming Tong , Hyoukjun Kwon, Tushar Krishna and Saibal Mukhopadhyay.
ISCA 2024 +LayoutLoop Integrated into NVLabs/timeloop	FEATHER: A Reconfigurable Accel. with Data Reordering Support for Low-Cost On-Chip Dataflow Switching Jianming Tong , Anirudh Itagi, Prasanth Chatarasi, Tushar Krishna International Symposium on Computer Architecture, Jun 2024 [Code] [Talk]
MLSys 2024	SmartPAF: Accurate Low-degree Polynomial Approximation of non-Polynomial Operators for Fast Private Inference in Homomorphic Encryption Jianming Tong *, Jingtian Dang*, Anupam Golder, Tushar Krishna. In Proc of Seventh Conference on Machine Learning and Systems, Jun 2024 [Code]
MLSys 2023 +Qualcomm Innovation Fellowship +Best Poster Award	SUSHI: SUbgraph Stationary Hardware-software Inference Co-design Payman Behnam*, Jianming Tong *, Alind, Yangyu, Yue, Pranav, Abhimanyu, Tushar, Alexey Tumanov In Proc of Sixth Conference on Machine Learning and Systems, Jun 2023
IEEE Micro 2023	Hardware-Software Co-design for Real-time Latency-Accuracy Navigation in TinyML Payman Behnam*, Jianming Tong *, Alind, Yangyu, Yue, Pranav, Abhimanyu, Tushar, Alexey Tumanov. IEEE Micro, Sep 2023
TRadar 2024	Real-time Digital RF Emulation – II: A Near Memory Custom Accelerator Xiangyu Mao, Mandovi Mukherjee, Nael Mianur Rahman, Coleman B DeLude, Joseph W. Driscoll, Sudarshan Sharma, Payman Behnam, Uday Kamal, Jongseok Woo, Daehyun Kim, Sharjeel M. Khan, Jianming Tong , Jamin Seo, Prachi Sinha, Madhavan Swaminathan, Tushar Krishna, Santosh Pande, Justin Romberg, and Saibal Mukhopadhyay. IEEE Transactions on Radar Systems, Sep 2024.

- TVLSI 2023** On Continuing DNN Accelerator Arch. Scaling Using Tightly-coupled Compute-on-Memory 3D ICs
 Gauthaman Murali, Aditya Iyer, Lingjun Zhu, **Jianming Tong**, Francisco Munoz Martinez, Srivatsa Rangachar Srinivasa, Tanay Karnik, Tushar Krishna, Sung Kyu Lim
 IEEE Transactions on Very Large Scale Integration Systems, Jul 2023
- RadarConf 2023** A High-Performance Computing Architecture for Real-Time Digital Emulation of RF Interactions
 Mandovi Mukherjee*, Nael Mizzanur Rahman*, Coleman B. DeLude*, Joseph W. Driscoll*, Uday Kamal, Jongseok Woo, Jamin Seo, Sudarshan Sharma, Xiangyu Mao, Payman Behnam, Sharjeel M. Khan, Daehyun Kim, **Jianming Tong**, Prachi Sinha, Santosh Pande, Tushar Krishna, Justin Romberg, Madhavan Swaminathan, and Saibal Mukhopadhyay.
 In Proc of IEEE Radar Conference, May 2023
- SENSORS 2023** SNATCH: Stealing Neural Network Architecture from ML Accelerator in Intelligent Sensors
 Sudarshan Sharma, Uday Kamal, **Jianming Tong**, Tushar Krishna, and Saibal Mukhopadhyay.
 IEEE SENSORS conference, Aug 2023.
- IMS 2023** FPGA-based High-Perf. Real-Time Emulation of Radar System using Direct Path Compute Model
 Xiangyu Mao*, Mandovi Mukherjee*, Nael M. Rahman*, Uday Kamal, Sudarshan Sharma, Payman Behnam, **Jianming Tong**, Jongseok Woo, Coleman B DeLude, Joseph W. Driscoll, Jamin Seo, Santosh Pande, Tushar Krishna, Justin Romberg, Madhavan Swaminathan, and Saibal Mukhopadhyay.
 In Proc of IEEE MTT-S International Microwave Symposium, Jun 2023
- IMS 2021** A Configurable Arch. for Efficient Sparse FIR Computation in Real-time Radio Frequency Systems
 Jamin Seo, Nael Mizzanur Rahman, Mandovi Mukherjee, Coleman DeLude, **Jianming Tong**, Justin Romberg, Tushar Krishna, and Saibal Mukhopadhyay.
 IEEE Microwave and Wireless Technology Letters, Sep 2022
- TOC 2021** PIT: Processing-In-Transmission with Fine-Grained Data Manipulation Networks
 Tian Xia, Pengchen Zong, Haoran Zhao, **Jianming Tong**, Wenzhe Zhao, Nanning Zheng and Pengju Ren.
 IEEE Transactions on Computers, Jul 2021
- FPT 2021** ac2SLAM: FPGA Accelerated High-Accuracy SLAM with Heapsort and Parallel Keypoint Extractor
 Cheng Wang, Yinkun Liu, Kedai Zuo, **Jianming Tong**, Yan Ding, and Pengju Ren.
 International Conference on Field-Programmable Technology, Jul 2021 [[Code](#)]
- ICRA 2021** SMMR-Explore: SubMap-based Multi-Robot Exploration System with Multi-robot Multi-target Potential Field Exploration Method
 Jincheng Yu*, **Jianming Tong***, Yuanfan Xu, Zhilin Xu, Haolin Dong, Tianxiang Yang and Yu Wang.
 IEEE International Conference on Robotics and Automation, Jan 2022 [[Code](#)][[Demo](#)]
- GLSVLSI 2020** Content-Oriented Configurable Architecture based on Highly Adaptive Data Transmission Networks
 Tian Xia, Pengchen Zong, Haoran Zhao, **Jianming Tong**, Wenzhe Zhao, Nanning Zheng and Pengju Ren.
 The 30th edition of the ACM Great Lakes Symposium on VLSI, Mar 2020

- ACS-DNN 2022** FastSwitch: Enabling Real-time DNN Switching via Weight-Sharing
Jianming Tong, Yangyu Chen, Yue Pan, Abhimanyu Bambhaniya, Alind Khare, Taekyung Heo, Alexey Tumanov, and Tushar Krishna
 Architecture, Compiler, and System Support for Multi-model DNN Workloads @ ISCA

PUBLICATIONS (* EQUAL CONTRIBUTION) -- BOOK

On Chip Networks, Second Edition [Translated Book in Mandarin]

Natalie Enright Jerger, Tushar Krishna, and Li-Shiuan Peh

Translator: Pengju Ren, Tian Xia, **Jianming Tong**, Pengcheng Zong, Haoran Zhao.

Publishing House of Electronics Industry, Jan 2021 [[Link](#)][[OriginalVersion](#)]

SELECTED TALKS

Leveraging ASIC AI chips for Homomorphic Encryption

Google Host: [Jeremy Kun](#) (May'24)

IBM Host: [Manoj Kumar](#), [Pradip Bose](#) (Aug'24)

NYU Host: [Brandon Reagon](#) (Nov'24)

UMich Host: [Todd Austin](#) (May'25)

Enabling Reconfigurable (Dataflow, Layout) CoSwitching in AI Accelerator

UT AUSTIN Host: [Mattan Erez](#) (Sep'25)

MIT Host: [Vivienne Sze](#) (Jun'24)

NVIDIA Host: [Angshuman Parashar](#) (Jul'24)

GaTech Host: [Alexandros Daglis](#) (Aug'24)

Enabling Real-time Accuracy Latency Navigation in Multi-Query AI Inference

HAN Lab @ MIT Host: [Hanrui Wang](#), [Song Han](#) (Oct'23)

EIC Lab@GaTech Host: [Celine Lin](#) (Jul'23)

A Sparse and Irregular GEMM Accelerator with Flexible Interconnects

Tsinghua Host: [Shulin Zeng](#) (Nov'22)

SELECTED PROJECTS

PROVE -- Leveraging ASIC AI chips for Zero-Knowledge Proof (ZKP)

Dec'24 – now

- Proposed Lazy Reduction, achieving up to an 18.5x speedup over high-precision Barrett Reduction.
- Proposed systematical memory intrinsics to explicitly explore dataflow factoring memory overhead.
- Deployed in Google cloud, achieving 111x end-to-end Multi-scalar Multiplication speedup over CPUs, enabling immediate Zero-knowledge Proof acceleration on commodity AI accelerators like TPUs [[link](#)].

CROSS -- Leveraging ASIC AI chips for Homomorphic Encryption (HE)

Jan'23 – now

- Proposed Basis Aligned Transformation to convert high-precision multiplication as low-precision MatMul.
- Proposed Layout Invariant 3-step Number Theory Transformation to convert NTT as Matrix Multiplication.
- CROSS achieves immediate privacy-preserving ML via TPUs (**126x, 5x faster than CPU, GPU-A100**).
- Integrated into Google privacy-preserving library and deployed in Google jaxite [[link](#)].

Enabling Privacy-preserving Real-time Multi-User VR Through Secure Outsourcing

Jul'23 – Jan'25

- Proposed **Horizontal Partitioning (HP)** to split multi-user VR flow into local-cloud for offloading less private data to untrusted cloud with noisy perturbation privacy protection for supporting more users.
- Developed, **Privatar**, the first framework leverages both local and untrusted cloud to concurrently achieve privacy-preserving multi-user VR, with **1.5x~2.27x** higher accuracy and **3.75x** more users support than the SotA completely model outsourcing. Such benefits only come at a negligible **9%** energy consumption.
- Applied PAC Privacy to multi-user VR, reducing noise intensity by up to **158x** compared to state-of-the-art differential privacy, achieving stronger privacy with minimal accuracy loss.

Approximating Non-linear Layers in ML Models for Homomorphic Encryption

Dec'22 – Jan'24

- Aim at reducing polynomial approximation degree of non-linear ML layers while preserving accuracy.
- Proposed, **SmartPAF**, the **first** training framework to replace non-linear operators with low-degree Polynomial Approximation Function and recover accuracy via ML fine-tuning, achieving **7.81x** speedup.
- Published [MLSys'24](#) with open-sourced [code](#), WIP to be integrated in Google [HEIR](#) compiler.

End-to-end Reconfigurable Flexible Machine Learning Accelerator (FEATHER, ISCA'24) Jan'21 – Jan'24

- Spotted layout switching as a performance-critical but often ignored issue in reconfigurable accelerators. A discordant layout slows down the theoretical performance of dataflow by up-to **120x**.
- Proposed **FEATHER**, the first architecture enabling (dataflow, layout) coswitching via novel NoC, **BIRRD**.
- Proposed **functional arbitrary reordering** to enable arbitrary layout switching and **implementational reordering in reduction** to hide layout reordering latency behind critical path.
- **Deployed on real FPGA**, achieving **2.65~4.56x** end-to-end throughput/PE improvement over SotAs.
- Published [ISCA'24](#) with open-sourced [code](#), layout modeling deployed in NVIDIA [Timeloop](#) library.

Enable Real-time Latency/Accuracy Navigation in Edge Applications

Mar'21 – Oct'22

- Worked on scheduling and hardware of multi-query inference system to improve performance.
- Proposed **SubGraph Stationary** to reuse shared weights of weight-shared networks across queries.
- Designed **SUSHI**, a multi-query inference serving system enabling SubGraph Stationary with novel hardware (SushiAccel) and software (SushiSched), improving **latency / accuracy** by **25% / 0.98%**.
- Published [MLSys'23](#), [IEEE Micro'23](#), wins [Qualcomm Innovation Fellowship](#), [Best Poster Award](#)@IAP'23.

Scalable Arbitrary Unicasting and Multicasting On-chip Network

Jan'21 – May'22

- Designed a scalable multi-stage on-chip network for **arbitrary multicasting** across hundred or **thousand nodes**, achieving **$O(N \log N)$ scalability** with the number of nodes N , better than $O(N^2)$ of crossbar.
- **Taped out** a realistic test chip for with 16 nodes under TSMC 28nm, verified on a real FPGA prototype.
- Published [TRadar'24](#), [IMs'23](#), [IMs'21](#), [RadarConf'23](#).

SKILLS

Programming (System) Verilog, Xilinx HLS, C/C++, Python, OpenCL, LLVM, MLIR, Clang
Tools Xilinx Vivado, Vitis HLS (AI), Cadence, Synapse

SERVICES

Reviewers ICRA'24, IROS'24, MLSys'25, CAL'25, IEEE Micro'25, TVLSI'25, ToC'25

AEC ASPLOS'23, ASPLOS'24, ISCA'24

Steering Team Computer Architecture Student Association ([CASA](#))

MEDIA COVERAGE

Semiconductor Engineering My paper “**Leveraging ASIC AI Chips for Homomorphic Encryption**” is added into Semiconductor Engineering’s library for its potential impact to boarder industry.
ACE News Jianming Tong: Spotlight from DARPA SRC JUMP2.0 Program ACE center (Aug'24)
GaTech News Jianming Tong: Ph.D. Students Named Rising Stars in Machine Learning (Jul'24)
GaTech News Jianming Tong: Ph.D. Students Won Qualcomm Innovation Fellowship (Jul'23)
GaTech News Jianming Tong won 2nd-place in SCS Poster Competition (Apr'23)

MENTORSHIP

	Name	First employment
MS MIT 2025	Nathan Xiong	MIT Master now
MS MIT 2025	Jan Strzeszynski	MIT Master now
MS GT 2024	Anirudh Itagi	Microsoft Azure AI Infrastructure
MS GT 2023	Yangyu Chen	Apple ASIC Verification Designer
UG GT 2022	Yue Pan	UCSD Ph.D.
UG GT 2022	Yuqi He	Apple ASIC Designer
UG GT 2022	Jingtian Dang	CMU ECE Master -> Now Ph.D. at GaTech
UG XJTU 2021	Yingkun Liu	SJTU Ph.D.
UG XJTU 2021	Kedai Zuo	UCSD Master
UG XJTU 2021	Cheng Wang	Tsinghua-XJTU Ph.D.