Jianming Tong

+1 470-357-8082 | jianming.tong@gatech.edu | jianmingtong.github.io

RESEARCH INTEREST

Computer Architecture with major interest on full-stack optimizations—spanning software, systems, and hardware—aimed at enhancing both the efficiency and privacy of AI systems.

EDUCATION

Jan'21 – May'26	**Georgia Institute of Technology, PhD** in Computer Science
	Advisor: Tushar Krishna, Focus: Computer Architecture
Jan'21 – May'24	**Georgia Institute of Technology**, **MS** in Computer Science
	Advisor: Tushar Krishna, Focus: Reconfigurable Dataflow Accelerator [ISCA'24]
Sep'16 – May'20	**Xi'an JiaoTong University**, **BE** in Electrical Engineering
	Advisor: Pengju Ren, *Thesis:* FPGA Accelerated High-Accuracy SLAM [FPT'21)

PROFESSIONAL EXPERIENCE

Aug'24 – Jan'25	**Student Researcher** in **Google**, Cambridge, MA
	Leveraging Google TPU for Homomorphic Encryption, host: **Asra Ali**, **Jevin Jiang**
Sep'23 – present	**Visiting Researcher** in **Massachusetts Institute of Technology**, Cambridge, MA
	Focus on Algorithmic, System and hardware acceleration for FHE, host: **Arvind**.
Jan'24 – May'1	**Teaching Assistant** in **Massachusetts Institute of Technology**, Cambridge, MA
	Constructive Computer Architecture (6.192) – Instructor: **Arvind**, **Tushar Krishna**
Jan'21 – present	**Graduate Research Assistant** in **Georgia Institute of Technology**, Atlanta, GA
	Software-System-Hardware Codesign for Edge ML (MLSys'23, IEEE MICRO'23)
May'23 – Aug' 23	**Engineer Intern** in **Rivos Inc.** Mountain View, CA, Mentor: **Gautham Chinya**
	Inter-chiplet performance modelling (NDA)
Jan'23 – May'23	**Teaching Assistant** in **Georgia Institute of Technology**, Atlanta, GA
	Processor Design (CS 3220) – Homework / Project Design
Jul' 22 – Aug' 22	**Research Intern** in **Pacific Northwest National Laboratory**, Remote
	A single-author end-to-end FPGA framework for AI inference (Tutorial@ICS'22)
Jun'21 – Aug'21	**Research Intern** in **DAMO Academy Alibaba Inc.** China, Mentor: **Jiansong Zhang**
	Designed and implemented FPGA accelerator for Homomorphic Encryption scalar/vector multiplication and addition (open-sourced with paper DAC'23)
Aug'20 – Jan'21	**Visiting Ph.D.** in **Tsinghua University**, Beijing, China, Mentor: **Yu Wang**
	Designed multi-robot collaborative exploration alg. (co-first author paper ICRA'21)
Sep'18 – Jul'20	**Undergraduate Research Assistant** in **Xi'an JiaoTong University**, Xi'an, China
	Lead NoC book translation and robotic HW acceleration (open-sourced paper FPT'21)

SELECTED AWARDS AND HONORS [Full list]

Jun'24	**ML and System Rising Star** @ MLCommon
Apr'24	**DAC Young Fellow** @ Design Automation Conference (DAC'24)
Sep'23	**Best Poster Award** – SUSHI @ Industry-Academia Partner Workshop (IAP'23)
	The top voted poster among over 20+ candidates by industry partners.
Jul'23	**Qualcomm Innovation Fellowship** – Latency/Accuracy Navigation in Edge ML
	18 winners out of 182 submissions (three rounds, nationwide)
Oct'22	**Qualcomm Innovation Fellowship Finalist** – ML Accel. Side-channel Attack

PUBLICATIONS (* EQUAL CONTRIBUTION) -- CONFERENCE

ISCA 2024	FEATHER: A Reconfigurable Accel. with Data Reordering Support for Low-Cost On-Chip Dataflow Switching
	Jianming Tong, Anirudh Itagi, Prasanth Chatarasi, Tushar Krishna.
	International Symposium on Computer Architecture, Jun 2024 [Code][Talk]

MLSys 2024	SmartPAF: Accurate Low-degree Polynomial Approximation of non-Polynomial Operators for Fast Private Inference in Homomorphic Encryption **Jianming Tong***, Jingtian Dang*, Anupam Golder, Tushar Krishna. In Proc of Seventh Conference on Machine Learning and Systems, Jun 2024 [Code]
MLSys 2023 +Qualcomm Innovation Fellowship +Best Poster Award	SUSHI: SUbgraph Stationary Hardware-software Inference Co-design Payman Behnam*, **Jianming Tong***, Alind, Yangyu, Yue, Pranav, Abhimanyu, Tushar, Alexey Tumanov. In Proc of Sixth Conference on Machine Learning and Systems, Jun 2023
IEEE Micro 2023	Hardware-Software Co-design for Real-time Latency-Accuracy Navigation in TinyML Payman Behnam*, **Jianming Tong***, Alind, Yangyu, Yue, Pranav, Abhimanyu, Tushar, Alexey Tumanov. IEEE Micro, Sep 2023
TRadar 2024	Real-time Digital RF Emulation – II: A Near Memory Custom Accelerator Xiangyu Mao, Mandovi Mukherjee, Nael Mizanur Rahman, Coleman B DeLude, Joseph W. Driscoll, Sudarshan Sharma, Payman Behnam, Uday Kamal, Jongseok Woo, Daehyun Kim, Sharjeel M. Khan, **Jianming Tong**, Jamin Seo, Prachi Sinha, Madhavan Swaminathan, Tushar Krishna, Santosh Pande, Justin Romberg, and Saibal Mukhopadhyay. IEEE Transactions on Radar Systems, Sep 2024.
TVLSI 2023	On Continuing DNN Accelerator Arch. Scaling Using Tightly-coupled Compute-on-Memory 3D ICs Gauthaman Murali, Aditya Iyer, Lingjun Zhu, **Jianming Tong**, Francisco Munoz Martinez, Srivatsa Rangachar Srinivasa, Tanay Karnik, Tushar Krishna,Sung Kyu Lim IEEE Transactions on Very Large Scale Integration Systems, Jul 2023
RadarConf 2023	A High-Performance Computing Architecture for Real-Time Digital Emulation of RF Interactions Mandovi Mukherjee*, Nael Mizanur Rahman*, Coleman B. DeLude*, Joseph W. Driscoll*, Uday Kamal, Jongseok Woo, Jamin Seo, Sudarshan Sharma, Xiangyu Mao, Payman Behnam, Sharjeel M. Khan, Daehyun Kim, **Jianming Tong**, Prachi Sinha, Santosh Pande, Tushar Krishna, Justin Romberg, Madhavan Swaminathan, and Saibal Mukhopadhyay. In Proc of IEEE Radar Conference, May 2023
SENSORS 2023	SNATCH: Stealing Neural Network Architecture from ML Accelerator in Intelligent Sensors Sudarshan Sharma, Uday Kamal, **Jianming Tong**, Tushar Krishna, and Saibal Mukhopadhyay. IEEE SENSORS conference, Aug 2023.
IMS 2023	FPGA-based High-Perf. Real-Time Emulation of Radar System using Direct Path Compute Model Xiangyu Mao*, Mandovi Mukherjee*, Nael M. Rahman*, Uday Kamal, Sudarshan Sharma, Payman Behnam, **Jianming Tong**, Jongseok Woo, Coleman B DeLude, Joseph W. Driscoll, Jamin Seo, Santosh Pande, Tushar Krishna, Justin Romberg, Madhavan Swaminathan, and Saibal Mukhopadhyay. In Proc of IEEE MTT-S International Microwave Symposium, Jun 2023
IMS 2021	A Configurable Arch. for Efficient Sparse FIR Computation in Real-time Radio Frequency Systems Jamin Seo, Nael Mizanur Rahman, Mandovi Mukherjee, Coleman DeLude, **Jianming Tong**, Justin Romberg, Tushar Krishna, and Saibal Mukhopadhyay. IEEE Microwave and Wireless Technology Letters, Sep 2022
TOC 2021	PIT: Processing-In-Transmission with Fine-Grained Data Manipulation Networks Tian Xia, Pengchen Zong, Haoran Zhao, **Jianming Tong**, Wenzhe Zhao, Nanning Zheng and Pengju Ren.

	IEEE Transactions on Computers, Jul 2021
FPT 2021	ac2SLAM: FPGA Accelerated High-Accuracy SLAM with Heapsort and Parallel Keypoint Extractor Cheng Wang, Yinkun Liu, Kedai Zuo, **Jianming Tong**, Yan Ding, and Pengju Ren. International Conference on Field-Programmable Technology, Jul 2021 [Code]
ICRA 2021	SMMR-Explore: SubMap-based Multi-Robot Exploration System with Multi-robot Multi-target Potential Field Exploration Method Jincheng Yu∗, **Jianming Tong∗**, Yuanfan Xu, Zhilin Xu, Haolin Dong, Tianxiang Yang and Yu Wang. IEEE International Conference on Robotics and Automation, Jan 2022 [Code][Demo]
GLSVLSI 2020	Content-Oriented Configurable Architecture based on Highly Adaptive Data Transmission Networks Tian Xia, Pengchen Zong, Haoran Zhao, **Jianming Tong**, Wenzhe Zhao, Nanning Zheng and Pengju Ren. The 30th edition of the ACM Great Lakes Symposium on VLSI, Mar 2020
ACS-DNN 2022	FastSwitch: Enabling Real-time DNN Switching via Weight-Sharing **Jianming Tong**, Yangyu Chen, Yue Pan, Abhimanyu Bambhaniya, Alind Khare, Taekyung Heo, Alexey Tumanov, and Tushar Krishna The 2nd Architecture, Compiler, and System Support for Multi-model DNN Workloads Workshop @ ISCA

PUBLICATIONS (* EQUAL CONTRIBUTION) -- BOOK

On Chip Networks, Second Edition [Translated Book in Mandarin]
Natalie Enright Jerger, Tushar Krishna, and Li-Shiuan Peh
Translator: Pengju Ren, Tian Xia, **Jianming Tong**, Pengcheng Zong, Haoran Zhao.
Publishing House of Electronics Industry, Jan 2021 [Link][OriginalVersion]

SELECTED TALKS

	Leveraging ASIC AI chips for Homomorphic Encryption
Google	Host: Jeremy Kun (May'24)
IBM	Host: Manoj Kumar, Pradip Bose (Aug'24)
NYU	Host: Brandon Reagon (Nov'24)
	Enabling Reconfigurable (Dataflow, Layout) CoSwitching in AI Accelerator
MIT	Host: Vivienne Sze (Jun'24)
NVIDIA	Host: Angshuman Parashar (Jul'24)
GaTech	Host: Alexandros Daglis (Aug'24)
	Enabling Real-time Accuracy Latency Navigation in Multi-Query AI Inference
HAN Lab @ MIT	Host: Hanrui Wang, Song Han (Oct'23)
EIC Lab@GaTech	Host: Celine Lin (Jul'23)
	A Sparse and Irregular GEMM Accelerator with Flexible Interconnects
Tsinghua	Host: Shulin Zeng (Nov'22)

SELECTED PROJECTS

Leveraging ASIC AI chips for Zero-Know Proof (ZKP)   Dec'24 – now
- WIP - Aiming at leveraging Google TPU for Zero-Knowledge Applications with better service price.

Leveraging ASIC AI chips for Homomorphic Encryption (HE)   Jan'23 – now
- Enabled immediate privacy-preserving ML by leveraging Google TPU to accelerate HE-based model.
- Proposed CROSS compiler to map high-precision polynomial multiplication on **TPUv4** as low-precision matrix multiplications and convolutions for acceleration on TPU (**126×, 5× faster than CPU, GPU-V100**).
- Integrated into Google privacy-preserving library and deployed in Google jaxite [link].

Enabling Privacy-preserving Real-time Multi-User VR Through Secure Outsourcing   Jul'23 – Jan'25
- Proposed **Horizontal Partitioning (HP)** to split multi-user VR flow into local-cloud for offloading less private data to untrusted cloud with noisy perturbation privacy protection for supporting more users.
- Developed, **Privatar**, the first framework leverages both local and untrusted cloud to concurrently achieve privacy-preserving multi-user VR, with **1.5×~2.27×** higher accuracy and **3.75×** more users support than the SotA completely model outsourcing. Such benefits only come at a negligible **9%** energy consumption.
- Applied PAC Privacy to multi-user VR, reducing noise intensity by up to **158×** compared to state-of-the-art differential privacy, achieving stronger privacy with minimal accuracy loss.

Approximating Non-linear Layers in ML Models for Homomorphic Encryption   Dec'22 – Jan'24
- Aim at reducing polynomial approximation degree of non-linear ML layers while preserving accuracy.
- Proposed, **SmartPAF**, the **first** training framework to replace non-linear operators with low-degree Polynomial Approximation Function and recover accuracy via ML fine-tuning, achieving **7.81×** speedup.
- Published [MLSys'24](#) with open-sourced [code](#), WIP to be integrated in Google [HEIR](#) compiler.

End-to-end Reconfigurable Flexible Machine Learning Accelerator (FEATHER, ISCA'24) Jan'21 – Jan'24
- Spotted layout switching as a performance-critical but often ignored issue in reconfigurable accelerators. A discordant layout slows down the theoretical performance of dataflow by up-to **120×**.
- Proposed **FEATHER**, the first architecture enabling (dataflow, layout) coswitching via novel NoC, **BIRRD**.
- Proposed **functional arbitrary reordering** to enable arbitrary layout switching and **implementational reordering in reduction** to hide layout reordering latency behind critical path.
- **Deployed on real FPGA**, achieving **2.65~4.56×** end-to-end throughput/PE improvement over SotAs.
- Published [ISCA'24](#) with open-sourced [code](#), layout modeling deployed in NVIDIA [Timeloop](#) library.

Enable Real-time Latency/Accuracy Navigation in Edge Applications   Mar'21 – Oct'22
- Worked on scheduling and hardware of multi-query inference system to improve performance.
- Proposed **SubGraph Stationary** to reuse shared weights of weight-shared networks across queries.
- Designed **SUSHI**, a multi-query inference serving system enabling SubGraph Stationary with novel hardware (SushiAccel) and software (SushiSched), improving **latency / accuracy** by **25% / 0.98%.**
- Published [MLSys'23](#), [IEEE Micro'23](#), wins [Qualcomm Innovation Fellowship](#), [Best Poster Award](#)@IAP'23.

Scalable Arbitrary Unicasting and Multicasting On-chip Network   Jan'21 – May'22
- Designed a scalable multi-stage on-chip network for **arbitrary multicasting** across hundred or **thousand nodes**, **achieving $O(NlogN)$ scalability** with the number of nodes N, better than $O(N^2)$ of crossbar.
- **Taped out** a realistic test chip for with 16 nodes under TSMC 28nm, verified on a real FPGA prototype.
- Published [TRadar'24](#), [IMS'23](#), [IMS'21](#), [RadarConf'23](#).

SKILLS
_____

Programming	(System) Verilog, Xilinx HLS, C/C++, Python, OpenCL, LLVM, MLIR, Clang
Tools	Xilinx Vivado, Vitis HLS (AI), Cadence, Synapse

SERVICES
_____

Reviewers	ICRA'24, IROS'24, MLSys'25
AEC	ASPLOS'23, ASPLOS'24, ISCA'24
Steering Team	Computer Architecture Student Association ([CASA](#))

MEDIA COVERAGE
_____

ACE News	Jianming Tong: Spotlight from DARPA SRC JUMP2.0 Program ACE center (Aug'24)
GaTech News	Jianming Tong: Ph.D. Students Named Rising Stars in Machine Learning (Jul'24)
GaTech News	Jianming Tong: Ph.D. Students Won Qualcomm Innovation Fellowship (Jul'23)
GaTech News	Jianming Tong won 2nd-place in SCS Poster Competition (Apr'23)

MENTORSHIP

	Name	First employment
MS MIT 2024	Jan Strzeszynski	MIT Master now
MS GT 2024	Anirudh Itagi	Microsoft Azure AI Infrastructure
MS GT 2023	Yangyu Chen	Apple ASIC Verification Designer
UG GT 2022	Yue Pan	UCSD Ph.D.
UG GT 2022	Yuqi He	Apple ASIC Designer
UG GT 2022	Jingtian Dang	CMU ECE Master -> Now Ph.D. at GaTech
UG XJTU 2021	Yingkun Liu	SJTU Ph.D.
UG XJTU 2021	Kedai Zuo	UCSD Master
UG XJTU 2021	Cheng Wang	Tsinghua-XJTU Ph.D.